

A Tour of Ranges in C++

Tom Peters, June 18, 2019

thomas.d.peters@gmail.com

tdp2110 on slack, github

@__tom_peters__

Outline

1. Introduction to ranges
2. Views
3. Actions
4. Algorithms
5. The dark side

range-v3

<https://github.com/ericniebler/range-v3/>

Work of Eric Niebler

- C++ Standards committee member
- Author of 4 boost libraries



What is a range?

What is a range? (Version 1)

A *range* is something that can be put in a range-based for-loop:

```
for (auto&& elt : rng) {  
    // do stuff with elt  
}
```

Examples:

- `std::vector`
- `std::array`
- `std::set`
- `std::string_view`

What is a range (Version 2)

A *range* is something we can call `std::begin` and `std::end`, which return *iterators*.

```
for (auto it = std::begin(rng); it != std::end(rng); ++it) {  
    // do stuff with it  
}
```

Iterators are generalizations of pointers, having:

- `operator*()`
- `operator++()`, `operator++(int)`

Iterators and the STL

```
std::vector<int> v = /* ... */;  
  
std::sort(v.begin(), v.end());  
auto it = std::find(v.begin(), v.end(),  
                    [](int elt) { return elt == 42; });  
auto it = std::maximum_element(v.begin(), v.end());
```



Range-based algorithms

// using range-v3 library

```
#include <range/v3/all.hpp>
```

```
std::vector<int> vec = /*...*/;
```

```
ranges::sort(vec);  
auto it = ranges::find(  
    vec,  
    predicate);
```

// C++20

```
#include <ranges>
```

```
std::vector<int> vec = /*...*/;
```

```
std::ranges::sort(vec)  
auto it = std::ranges::find(  
    vec,  
    predicate);
```


Composing STL algorithms is awkward

```
// given  
std::vector<int> v = {0,1,2,3,4,5,6,7,8,9};  
  
// how to produce a vector consisting of the squares  
// of the even elements?
```

Composing STL algorithms is awkward

```
#include <algorithm>
#include <vector>

std::vector<int> v = {1,2,3,4,5};

std::vector<int> evens;
std::copy_if(v.begin(),
            v.end(),
            std::back_inserter(evens),
            [](int elt) { return elt % 2 == 0; });

std::vector<int> res;
std::transform(evens.begin(),
              evens.end(),
              std::back_inserter(res),
              [](int elt) { return elt * elt; });
```

Range version:

```
#include <vector>
#include <range/v3/all.hpp>

std::vector<int> v = {1,2,3,4,5};

auto rng = v | ranges::view::filter(
    [](int elt) { return elt % 2 == 0; })
    | ranges::view::transform(
    [](int elt) { return elt * elt; });
```

Range concepts

```
template<class T>
concept bool Range =
    requires(T&& t) {
        { std::begin(t) } -> Iterator<T>
        { std::end(t) } -> Iterator<T>
    }
}
```

*// see
<https://github.com/CaseyCarter/cmcstl2/>
// for proper definitions*

```
template<class T>
concept bool View =
    Range<T> &&
    Semiregular<T> &&
    ViewPredicate<T>;
```

*// see
<https://en.cppreference.com/w/cpp/experimental/ranges/range/View>*

Range concepts

```
template<class T>
concept bool Range =
    requires(T&& t) {
        { std::begin(t) } -> Iterator<T>
        { std::end(t) } -> Iterator<T>
    }
}

template<class T>
concept bool View =
    Range<T> &&
    Semiregular<T> &&
    ViewPredicate<T>;
```

```
template<class T>
    requires Range<T>
void Foo(T&& t) {
    // ...
}
```

What is a range (Version 3)

```
template<class T>
concept bool Range =
    requires(T&& t) {
        { std::begin(t) } -> Iterator<T>
        { std::end(t) } -> Iterator<T>
    }
}
```

A *range* is something which models the Range concept.

A few range adaptors

Filter, transform

- **Filter.** operator++ is interesting: skips until predicate is met
- **Transform.** operator* is interesting: it applies a predicate (only when asked for)

Range adaptors: all

```
auto rng2 = rng1 | view::all;
```

Range adaptors: Concat

Join: traverse first range, then second range, then ...

```
std::vector<int> v = {1,2,3,4};  
std::set<int> s = {5,6,7,8,9};  
  
auto rng = ranges::view::concat(v, s);  
  
// traversing yields 1,2,3,4, 5,6,7,8,9
```

Range adaptors: take and slice

```
std::vector<int> v = {1,2,3,4,5,6,7,8,9,10};
```

```
auto first_five = v | ranges::view::take(5);
```

```
auto slice = v | ranges::view::slice(3,6);
```

Range Adaptors: Join

Flatten a range of ranges to a single range:

```
std::vector<std::vector<int>> v = {{1,2}, {3,4,5}, {6,7,8,9}};
```


```
std::vector<int> flat = v | ranges::view::join;
```

```
// flat = {1,2,3,4,5,6,7,8,9};
```

Zip

```
std::vector<int> v = {1,2,3,4};  
std::set<int> s = {5,6,7,8,9};  
  
auto rng = ranges::view::zip(v, s);  
  
// traversal gives pairs of references:  
// (1,5), (2,6), (3,7), (4,8)
```

<http://ericniebler.com/2015/01/28/to-be-or-not-to-be-an-iterator/>
<https://ericniebler.github.io/std/wg21/D0022.html>



Enumerate

// range-v3 version

```
std::set<std::string> s =  
    {"hello", "world"};  
  
for (auto const& [index, str] :  
    ranges::view::enumerate(s))  
{  
    // do stuff  
}
```

// old style

```
std::set<std::string> v =  
    {"hello", "world"};  
  
int index = 0;  
for (auto const& str : s) {  
    // do stuff  
    ++index;  
}
```

Range generators

```
auto nats = ranges::view::ints();
```

```
// yields 0,1,2,3, ...
```

```
// example, enumerate can be written:
```

```
auto enumerator = ranges::view::zip(ranges::view::ints(),  
someRng);
```

```
auto first_ten_ints = nats | ranges::view::take(10);
```

Range actions

Eager, mutating, composable algorithms

```
std::vector<T> v;
```

```
v = move(v) | action::sort | action::unique;
```


Range actions

<ul style="list-style-type: none"><code>adjacent_remove_if</code><code>drop</code><code>drop_while</code><code>erase</code><code>insert</code><code>join</code><code>remove</code><code>remove_if</code><code>reverse</code><code>shuffle</code>	<ul style="list-style-type: none"><code>slice</code><code>sort</code><code>split</code><code>split_when</code><code>stable_sort</code><code>stride</code><code>take</code><code>take_while</code><code>transform</code><code>unique</code>
---	---

Projections

```
struct Person {  
    int id;  
};  
std::vector<Person> v;
```

```
std::sort(v.begin(), v.end(), [](auto& p1, auto& p2) {  
    return p1.id < p2.id; });
```

```
ranges::sort(v, std::less{}, [](auto& p) { return p.id; });
```

```
ranges::sort(v, std::less{}, &Person::id);
```



Projections && `std::invoke`

```
struct Person {  
    int id;  
};
```

What is `&Person::id`??

Pointer to member data: type `int Person::*`

May be called dynamically:

```
Person p {42};  
int Person::* memberDataPtr = &Person::id;  
assert(p.id == p.*memberDataPtr);  
assert(p.id == std::invoke(memberDataPtr, p));  
assert(p.id == std::invoke([](Person& p) { return p.id; }, p));
```

Instead of `f(args...)`, range-v3 uses the more general `std::invoke(f, args...)`

But wait, there's more!

- Templates for generating views!
- Templates for generating adaptors!
- (Work in progress) asynchronous ranges!



```
dharam@dharam-H110MHC: ~  
dharam@dharam-H110MHC:~$ cal 2018  
2018  
January February March  
Su Mo Tu We Th Fr Sa Su Mo Tu We Th Fr Sa Su Mo Tu We Th Fr Sa  
1 2 3 4 5 6 1 2 3 1 2 3  
7 8 9 10 11 12 13 4 5 6 7 8 9 10 4 5 6 7 8 9 10  
14 15 16 17 18 19 20 11 12 13 14 15 16 17 11 12 13 14 15 16 17  
21 22 23 24 25 26 27 18 19 20 21 22 23 24 18 19 20 21 22 23 24  
28 29 30 31 25 26 27 28 25 26 27 28 29 30 31  
  
April May June  
Su Mo Tu We Th Fr Sa Su Mo Tu We Th Fr Sa Su Mo Tu We Th Fr Sa  
1 2 3 4 5 6 7 1 2 3 4 5 1 2  
8 9 10 11 12 13 14 6 7 8 9 10 11 12 3 4 5 6 7 8 9  
15 16 17 18 19 20 21 13 14 15 16 17 18 19 10 11 12 13 14 15 16  
22 23 24 25 26 27 28 20 21 22 23 24 25 26 17 18 19 20 21 22 23  
29 30 27 28 29 30 31 24 25 26 27 28 29 30  
  
July August September  
Su Mo Tu We Th Fr Sa Su Mo Tu We Th Fr Sa Su Mo Tu We Th Fr Sa  
1 2 3 4 5 6 7 1 2 3 4 1  
8 9 10 11 12 13 14 5 6 7 8 9 10 11 2 3 4 5 6 7 8  
15 16 17 18 19 20 21 12 13 14 15 16 17 18 9 10 11 12 13 14 15  
22 23 24 25 26 27 28 19 20 21 22 23 24 25 16 17 18 19 20 21 22  
29 30 31 26 27 28 29 30 31 23 24 25 26 27 28 29  
30  
  
October November December  
Su Mo Tu We Th Fr Sa Su Mo Tu We Th Fr Sa Su Mo Tu We Th Fr Sa  
1 2 3 4 5 6 1 2 3 1  
7 8 9 10 11 12 13 4 5 6 7 8 9 10 2 3 4 5 6 7 8  
14 15 16 17 18 19 20 11 12 13 14 15 16 17 9 10 11 12 13 14 15  
21 22 23 24 25 26 27 18 19 20 21 22 23 24 16 17 18 19 20 21 22  
28 29 30 31 25 26 27 28 29 30 23 24 25 26 27 28 29  
30 31
```



Compiler support

From <https://ericniebler.github.io/range-v3/>

- clang 3.6.2
- GCC 4.9.1
- MSVC VS2017 15.9 (`_MSC_VER >= 1916`), with `/std:c++17 /permissive-`

Requirements

- Requires C++11/14/17
- pre-C++17, can't use range-based for on all ranges, need **RANGES_FOR** macro.

Downsides of ranges

Ownership



Arvid Gerstmann

@ArvidGerstmann

Follow



C++ Community: We have this cool concept called RAI, which allows us to write code which has no dangling pointers, resource leaks or use-after-free

Eric Niebler: Hold my beer

Ownership 2

```
auto rng = std::vector<int>{1,2,3} | ranges::view::reverse;
```

error: static assertion failed: You can't pipe an rvalue container into a view. First, save the container into a named variable, and then pipe it to the view.

```
static_assert(ranges::View<Rng>() || std::is_lvalue_reference<Rng>(),
```


~~~~~^~~~~~

# Ownership 3

```
auto Oops()  
{  
    std::vector<int> v{1,2,3,4};  
    return v | ranges::view::reverse;  
}  
  
int main() {  
    for (auto elt : Oops()) { // use after free!  
        // whatever  
    }  
}
```

# ICE

fatal error C1001: An internal error has occurred in the compiler.



# A few more of my complaints:

1. No view to const
2. Creating custom ranges/views is tedious
3. Error messages can be bad.

# "Modern" C++ lamentations

<https://aras-p.info/blog/2018/12/28/Modern-C-Lamentations/>



# Pythagorean Triples

A triple of integers  $(a,b,c)$  such that  $a^2 + b^2 = c^2$ .

Examples:

- $(3, 4, 5): \quad 3^2 + 4^2 = 5^2$
- $(5, 12, 13): \quad 5^2 + 12^2 = 13^2$

# Pythagorean Triples

Printing the first n:

```
void printNTriples(int n)
{
    int i = 0;
    for (int z = 1; ; ++z)
        for (int x = 1; x <= z; ++x)
            for (int y = x; y <= z; ++y)
                if (x*x + y*y == z*z) {
                    printf("%d, %d, %d\n", x, y, z);
                    if (++i == n)
                        return;
                }
}
```



# Pythagorean triples, lazily

From <http://ericniebler.com/2018/12/05/standard-ranges/>

```
auto triples =  
    for_each(iota(1), [](int z) {  
        return for_each(iota(1, z+1), [=](int x) {  
            return for_each(iota(x, z+1), [=](int y) {  
                return yield_if(x*x + y*y == z*z,  
                    make_tuple(x, y, z));  
            });  
        });  
    });
```

# Pythagorean triples, lazily

From <http://ericniebler.com/2018/12/05/standard-ranges/>

```
auto triples =  
    for_each(iota(1), [](int z) {  
        return for_each(iota(1, z+1), [=](int x) {  
            return for_each(iota(x, z+1), [=](int y) {  
                return yield_if(x*x + y*y == z*z,  
                    make_tuple(x, y, z));  
            });  
        });  
    });
```

```
for_each(rng, f) <=> rng | transform(f) | join;
```

# Pythagorean triples, lazily

From <http://ericniebler.com/2018/12/05/standard-ranges/>

```
auto triples =  
  for_each(iota(1), [])(int z) {  
    return for_each(iota(1, z+1), [=](int x) {  
      return for_each(iota(x, z+1), [=](int y) {  
        return yield_if(x*x + y*y == z*z,  
          make_tuple(x, y, z));  
      });  
    });  
  });
```

```
-- Haskell  
triples :: [(Int, Int, Int)]  
triples = [(x, y, z)  
  | z <- [0..]  
  , x <- [1..z]  
  , y <- [x..z]  
  , x * x + y * y == z * z  
  ]
```

```
for_each(rng, f) = rng | transform(f) | join;
```

# Pythagorean triples

```
auto triples =  
    for_each(iota(1), [](int z) {  
        return for_each(iota(1, z+1), [=](int x) {  
            return for_each(iota(x, z+1), [=](int y) {  
                return yield_if(x*x + y*y == z*z,  
                    make_tuple(x, y, z));  
            });  
        });  
    });  
for (auto elt: triples | ranges::view::take(100)) {  
    std::cout << elt << "\n";  
}
```

Compile time: ~3 sec

Runtime (debug): 300ms (0.3s)

Runtime (release): 1ms

```
void printNTriples(int n)  
{  
    int i = 0;  
    for (int z = 1; ; ++z)  
        for (int x = 1; x <= z; ++x)  
            for (int y = x; y <= z; ++y)  
                if (x*x + y*y == z*z) {  
                    printf("%d, %d, %d\n", x, y, z);  
                    if (++i == n)  
                        return;  
                }  
}  
  
printNTriples(100);
```

0.064 sec

2ms

0ms

# Coroutines

```
#include <cppcoro/generator.hpp>
#include <tuple>

cppcoro::generator<std::tuple<int,int,int>> triples()
{
    for (int z = 1; ; ++z)
        for (int x = 1; x <= z; ++x)
            for (int y = x; y <= z; ++y)
                if (x*x + y*y == z*z) {
                    co_yield std::make_tuple(x,y,z);
                }
}
```

# Conclusions

- Ranges give new primitives to more directly state intent.
- But they come with a cost: compile times and debug perf.